

Machine Learning on Kubernetes

Creating and uploading necessary files in GCP- Cloud Shell Terminal

1. Start minikube in Google Cloud Platform

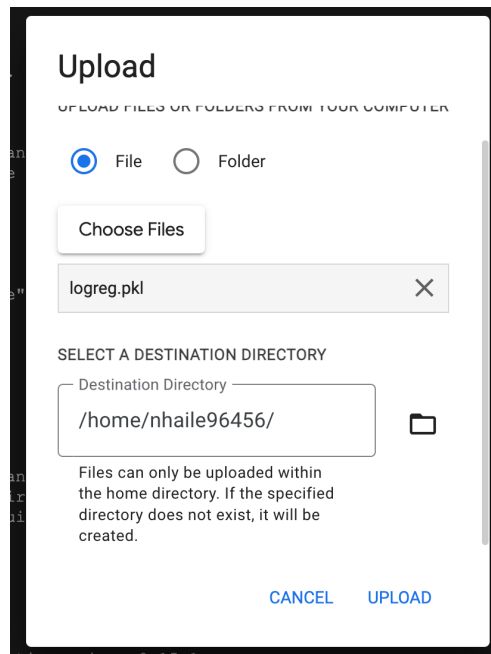
- *minikube start*

```
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$ minikube start
* minikube v1.33.1 on Ubuntu 22.04 (amd64)
  - MINIKUBE_FORCE_SYSTEMD=true
  - MINIKUBE_HOME=/google/minikube
  - MINIKUBE_WANTUPDATENOTIFICATION=false
* Using the docker driver based on existing profile
* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.44 ...
* Updating the running docker "minikube" container ...
* Preparing Kubernetes v1.30.0 on Docker 26.1.1 ...
  - kubelet.cgroups-per-qos=false
  - kubelet.enforce-node-allocatable=""
* Verifying Kubernetes components...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: storage-provisioner, default-storageclass
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$
```

2. Create a requirements.txt file.

```
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$ vi requirements.txt
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$ cat requirements.txt
Flask==1.1.1
gunicorn==19.9.0
itsdangerous==1.1.0
Jinja2==2.10.1
MarkupSafe==1.1.1
Werkzeug==0.15.5
numpy==1.19.5 # Adjusted to a version before np.float deprecation scipy>=0.15.1
scikit-learn==0.24.2 # Ensure compatibility with numpy version matplotlib>=1.4.3
pandas>=0.19
flasgger==0.9.4
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$
```

3. Upload logreg.pkl file to your workspace.



Upload

UPLOAD FILES OR FOLDERS FROM YOUR COMPUTER

☒ File ☐ Folder

Choose Files

logreg.pkl

SELECT A DESTINATION DIRECTORY

Destination Directory

/home/nhaile96456/

Files can only be uploaded within the home directory. If the specified directory does not exist, it will be created.

CANCEL UPLOAD

4. Create a flask_api.py file and paste the below python code in it.

```
from flask import Flask, request
import numpy as np
import pickle
import pandas as pd
from flasgger import Swagger

app = Flask(__name__)
Swagger(app)

# Load the logistic regression model
pickle_in = open("logreg.pkl", "rb")
model = pickle.load(pickle_in)

@app.route('/')
def home():
    return "Welcome to the Flask API!"

@app.route('/predict', methods=["GET"])
def predict_class():
    """
    Predict if Customer would buy the product or not.
    """
    parameters:
    - name: age
      in: query
      type: number
      required: true
    - name: new_user
      in: query
```

```

    type: number
    required: true
- name: total_pages_visited
  in: query
  type: number
  required: true
responses:
  200:
    description: Prediction
"""

age = int(request.args.get("age"))
new_user = int(request.args.get("new_user"))
total_pages_visited = int(request.args.get("total_pages_visited"))
prediction = model.predict([[age, new_user, total_pages_visited]])
return "Model prediction is " + str(prediction[0])

@app.route('/predict_file', methods=["POST"])
def prediction_test_file():
    """
    Prediction on multiple input test file.
    ---
    parameters:
      - name: file
        in: formData
        type: file
        required: true
    responses:
      200:
        description: Test file Prediction
    """
    df_test = pd.read_csv(request.files.get("file"))
    prediction = model.predict(df_test)
    return str(list(prediction))

if __name__ == '__main__':
    app.run(debug=True, host='0.0.0.0', port=5000)

```

```

nhaile96456@cloudshell:~ (cs570-big-data-analytics) $ vi flask_api.py
nhaile96456@cloudshell:~ (cs570-big-data-analytics) $ cat flask_api.py
from flask import Flask, request
import numpy as np
import pickle
import pandas as pd
from flasgger import Swagger

app = Flask(__name__)
Swagger(app)

# Load the logistic regression model
pickle_in = open("logreg.pkl", "rb")
model = pickle.load(pickle_in)

@app.route('/')
def home():
    return "Welcome to the Flask API!"

@app.route('/predict', methods=["GET"])
def predict_class():
    """

```

5. Create a DockerFile with the following content in it.

```
# Use the official Python image from the Docker Hub
FROM python:3.8-slim
# Set the working directory in the container
WORKDIR /app
# Copy the current directory contents into the container at /app
COPY . /app
# Install any needed packages specified in requirements.txt
RUN pip install --no-cache-dir -r requirements.txt
# Make port 5000 available to the world outside this container
EXPOSE 5000
# Define environment variable to prevent Python from writing .pyc files to disk
ENV PYTHONUNBUFFERED=1
# Run flask_api.py when the container launches
CMD ["python", "flask_api.py"]
```

```
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$ vi Dockerfile
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$ cat Dockerfile
# Use the official Python image from the Docker Hub
FROM python:3.8-slim

# Set the working directory in the container
WORKDIR /app

# Copy the current directory contents into the container at /app
COPY . /app

# Install any needed packages specified in requirements.txt
RUN pip install --no-cache-dir -r requirements.txt

# Make port 5000 available to the world outside this container
EXPOSE 5000

# Define environment variable to prevent Python from writing .pyc files to disk
ENV PYTHONUNBUFFERED=1

# Run flask_api.py when the container launches
CMD ["python", "flask_api.py"]

nhaile96456@cloudshell:~ (cs570-big-data-analytics)$
```

Explanation

FROM python:3.8-slim: This line specifies the base image for the Docker container, which is a lightweight version of Python 3.8.

Working Directory:

WORKDIR /app: This sets the working directory to /app in the container. All subsequent commands will be run from this directory.

Copy Files:

COPY . /app: This copies the contents of the current directory on your host machine to the /app directory in the container.

Install Dependencies:

RUN `pip install --no-cache-dir -r requirements.txt`: This installs the dependencies listed in `requirements.txt` using `pip`. The `--no-cache-dir` option reduces the size of the image by not caching the downloaded packages.

Expose Port:

EXPOSE 5000: This makes port 5000 available for network connections. It does not actually publish the port; it simply serves as documentation.

Environment Variable:

ENV PYTHONUNBUFFERED=1: This prevents Python from buffering stdout and stderr, ensuring that logs are immediately available in Docker.

Command:

CMD `["python", "flask_api.py"]`: This specifies the command to run the Flask application when the container starts.

6. To build the docker image use the command.

`sudo docker build -t ml_app_docker .`

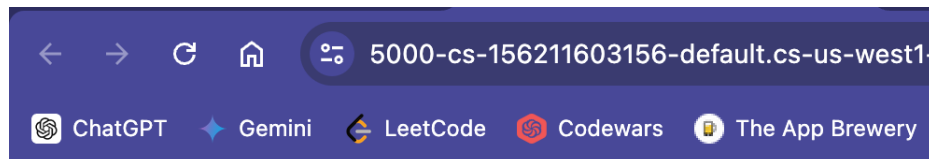
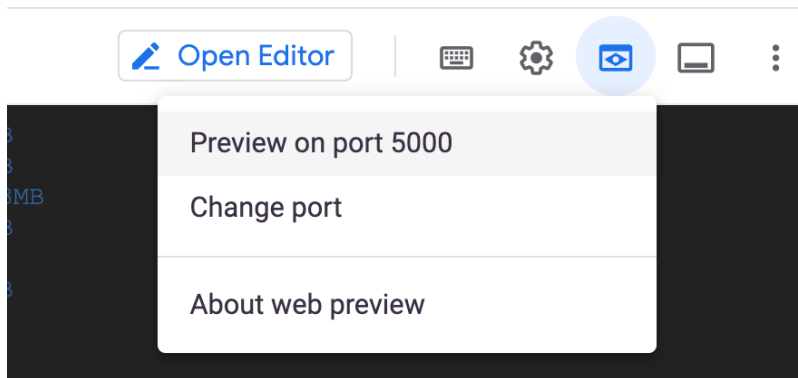
```
nhaille96456@cloudshell:~ (cs570-big-data-analytics)$ sudo docker build -t ml_app_docker .
[+] Building 39.0s (9/9) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 621B
=> [internal] load metadata for docker.io/library/python:3.8-slim
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [1/4] FROM docker.io/library/python:3.8-slim@sha256:c177b5b444d6913678d80bd26af131187de166cd68ac66605acbbc76elb343d7
=> => resolve docker.io/library/python:3.8-slim@sha256:c177b5b444d6913678d80bd26af131187de166cd68ac66605acbbc76elb343d7
=> => sha256:fea4f2170757fa9e87b6421086bf5d32880bd009555941641108ea89519b5742 11.67MB / 11.67MB
=> => sha256:c177b5b444d6913678d80bd26af131187de166cd68ac66605acbbc76elb343d7 10.41kB / 10.41kB
=> => sha256:b2d6266a63eafff4fce6a9149b80686ff6c17b0e4556d0320198ce85e9aa41d 1.94kB / 1.94kB
=> => sha256:5208c64e1783e15b2d0ee357a3979688ae3faa21c6a43f4965e3530d68abbcfc 6.93kB / 6.93kB
=> => sha256:efc2b5ad9ee00b5fa54239d53feaa3569ccbef689aa5e5dbfc25da6c4df559 29.13MB / 29.13MB
=> => sha256:0d9335f02ede5b557e3899b4161a3a7f7d8461fd62d558451b3884172a710962 3.51MB / 3.51MB
=> => sha256:e5635d0cdd4c514a4e76d97174785a452a1e81b87b6a8bd8ce09c5ac2d135df8 230B / 230B
=> => sha256:ebe530eb534fda723884e0b61fa4633bb792a9600f452a779301c8a7e4216d83 2.78MB / 2.78MB
=> => extracting sha256:efc2b5ad9ee00b5fa54239d53feaa3569ccbef689aa5e5dbfc25da6c4df559 2.2s
=> => extracting sha256:0d9335f02ede5b557e3899b4161a3a7f7d8461fd62d558451b3884172a710962 0.2s
=> => extracting sha256:fea4f2170757fa9e87b6421086bf5d32880bd009555941641108ea89519b5742 0.6s
=> => extracting sha256:e5635d0cdd4c514a4e76d97174785a452a1e81b87b6a8bd8ce09c5ac2d135df8 0.0s
=> => extracting sha256:ebe530eb534fda723884e0b61fa4633bb792a9600f452a779301c8a7e4216d83 0.3s
=> [internal] load build context
=> => transferring context: 62.48MB
=> [2/4] WORKDIR /app
=> [3/4] COPY . /app
=> [4/4] RUN pip install --no-cache-dir -r requirements.txt
=> exporting to image
=> exporting layers
=> writing image sha256:e0elb675538fcdc838c47bcc32929fdeaf6aeab562e2d22a848eac4207e97e4
=> naming to docker.io/library/ml_app_docker
nhaille96456@cloudshell:~ (cs570-big-data-analytics)$
```

7. The following command runs a Docker container from the `ml_app_docker` image.

`docker container run -p 5000:5000 ml_app_docker`

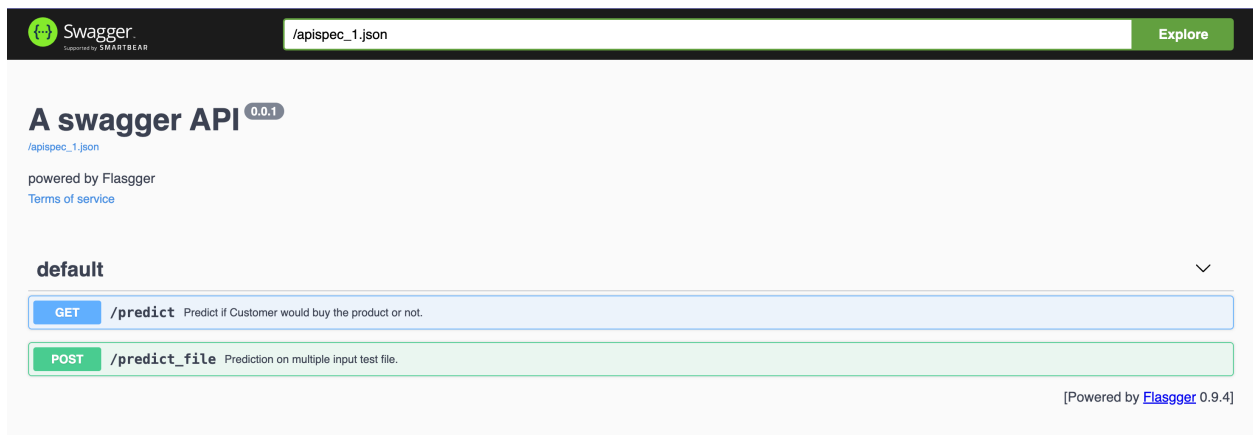
```
nhaille96456@cloudshell:~ (cs570-big-data-analytics)$ docker container run -p 5000:5000 ml_app_docker
/usr/local/lib/python3.8/site-packages/sklearn/base.py:310: UserWarning: Trying to unpickle estimator LogisticRegression from version 0.23.2 when using version 0.24.2.
This might lead to breaking code or invalid results. Use at your own risk.
warnings.warn(
* Serving Flask app "flask_api" (lazy loading)
* Environment: production
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Debug mode: on
* Running on http://0.0.0.0:5000/ (Press CTRL+C to quit)
* Restarting with stat
/usr/local/lib/python3.8/site-packages/sklearn/base.py:310: UserWarning: Trying to unpickle estimator LogisticRegression from version 0.23.2 when using version 0.24.2.
This might lead to breaking code or invalid results. Use at your own risk.
warnings.warn(
* Debugger is active!
* Debugger PIN: 269-833-583
```

- On the top right side of the screen, click the little eye shaped button. Change the port number if it's not already 5000. You will expect to see some sort of a welcome message.



Welcome to the Flask API!

- Add /apidocs/ at the end of the URL and you will see the home page of Swagger-UI.



10. Click ‘Get’ and ‘Try it out’ at the top right side to get the following page.

GET

/predict

Predict if Customer would buy the product or not.

Parameters

Cancel

Name	Description
age <small>required</small> number <i>(query)</i>	<input type="text" value="age"/>
new_user <small>required</small> number <i>(query)</i>	<input type="text" value="new_user"/>
total_pages_visited <small>required</small> number <i>(query)</i>	<input type="text" value="total_pages_visited"/>

Execute

Responses

Response content type application/json

Code	Description
200	Prediction

POST

/predict_file

Prediction on multiple input test file.

11. Enter values for the input parameters and click ‘Execute.’

GET

/predict

Predict if Customer would buy the product or not.

Parameters

Cancel

Name	Description
age <small>required</small> number <i>(query)</i>	<input type="text" value="23"/>
new_user <small>required</small> number <i>(query)</i>	<input type="text" value="2"/>
total_pages_visited <small>required</small> number <i>(query)</i>	<input type="text" value="5"/>

ExecuteClear

Responses

Response content type application/json

12. Upon the execution call, the request goes to the app and predictions are made by the model.

The result of the model prediction is displayed in the Prediction section of the page as following

The screenshot displays a REST client interface with the following sections:

- Responses:** A dropdown menu set to `application/json`.
- Curl:** A text box containing the command: `curl -X GET "https://5000-cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev/predict?age=23&new_user=2&total_pages_visited=5" -H "accept: application/json"`
- Request URL:** A text box containing: `https://5000-cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev/predict?age=23&new_user=2&total_pages_visited=5`
- Server response:**
 - Code:** 200
 - Details:**
 - Response body:** A text box showing `Model prediction is 0` with a `Download` button.
 - Response headers:** A text box showing:

```
content-length: 21
content-security-policy: frame-ancestors 'self' https://80-cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev https://cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev https://ide.cloud.google.com https://shell.cloud.google.com https://ssh.cloud.google.com https://console.cloud.google.com
content-type: text/html; charset=utf-8
date: Mon, 29 Jul 2024 07:47:21 GMT
server: Werkzeug/0.15.5 Python/3.8.19
```
- Responses:** A table with the following data:

Code	Description
200	Prediction

13. Next, the app can make predictions for a group of customers (test data) by clicking ‘Post’.

The screenshot shows the `POST /predict_file` endpoint interface. It includes a `Try it out` button, a **Parameters** section with a `file` input (labeled `file * required` and `file (formData)`), and a **Responses** section with a `Response content type` dropdown set to `application/json`. The response table shows a `200` status for `Test file Prediction`.

14. Upload the `test_data.csv` file and click ‘Execute.’

This screenshot shows the same `POST /predict_file` interface, but with the `file` input field containing the text `test_data.csv`. A blue `Execute` button is highlighted, and a red `Cancel` button is visible in the top right corner. The **Responses** section remains the same.

15. The model would make the predictions, and the results would be displayed as shown below.

Curl	
curl -X POST "https://5000-cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev/predict_file" -H "accept: application/json" -H "Content-Type: multipart/form-data" -F "file=test_data.csv; type=text/csv"	
Request URL	https://5000-cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev/predict_file
Server response	
Code	Details
200	<div><div>Response body</div><div>[0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0]</div><div>Download</div></div> <div><div>Response headers</div><div>access-control-allow-credentials: true access-control-allow-methods: GET, POST, OPTIONS, PATCH, DELETE access-control-allow-origin: https://5000-cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev content-length: 150 content-security-policy: frame-ancestors 'self' https://80-cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev https://cs-156211603156-default.cs-us-west1-ijlt.cloudshell.dev https://ide.cloud.google.com https://shell.cloud.google.com https://ssh.cloud.google.com https://console.cloud.google.com content-type: text/html; charset=utf-8 date: Mon, 29 Jul 2024 07:52:31 GMT server: Werkzeug/0.15.5 Python/3.8.19</div></div>
Responses	
Code	Description
200	Test file Prediction

16. To list running docker containers to kill them, run the command `docker ps`.

```

nhaile96456@cloudshell:~ (cs570-big-data-analytics)$ docker ps
CONTAINER ID   IMAGE                                COMMAND                                     CREATED        STATUS        PORTS
6393b5a57ffe9   gcr.io/k8s-minikube/kicbase:v0.0.44  "/usr/local/bin/entr..."  39 minutes ago   Up 39 minutes   127.0.0.1:32768->22/tcp, 127.0.0.1:32769->2376/tcp, 127.0.0.1:32770->5000/tcp, 127.0.0.1:32771->8443/tcp, 127.0.0.1:32772->32443/tcp
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$

```

17. Then run `'docker kill <container id>'`.

```
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$ docker kill 6393b5a57ff9
6393b5a57ff9
nhaile96456@cloudshell:~ (cs570-big-data-analytics)$
```